

```
In [7]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```


```
In [8]: df=pd.read_csv("Employee.csv")
```

```
In [9]: df.head()
```

Out[9]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeN
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns



In [10]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                            1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                      1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [11]: df.isnull().any()
```

```
Out[11]: Age                False
Attrition                  False
BusinessTravel             False
DailyRate                 False
Department                False
DistanceFromHome          False
Education                 False
EducationField             False
EmployeeCount             False
EmployeeNumber            False
EnvironmentSatisfaction   False
Gender                    False
HourlyRate                False
JobInvolvement            False
JobLevel                  False
JobRole                   False
JobSatisfaction           False
MaritalStatus             False
MonthlyIncome            False
MonthlyRate              False
NumCompaniesWorked        False
Over18                    False
OverTime                  False
PercentSalaryHike         False
PerformanceRating         False
RelationshipSatisfaction  False
StandardHours             False
StockOptionLevel          False
TotalWorkingYears         False
TrainingTimesLastYear     False
WorkLifeBalance           False
YearsAtCompany            False
YearsInCurrentRole        False
YearsSinceLastPromotion   False
YearsWithCurrManager      False
dtype: bool
```

In [12]: `df.corr()`

```
C:\Users\Praveen\AppData\Local\Temp\ipykernel_25940\1134722465.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.  
df.corr()
```

Out[12]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfacti
Age	1.000000	0.010661	-0.001686	0.208034	NaN	-0.010145	0.0101
DailyRate	0.010661	1.000000	-0.004985	-0.016806	NaN	-0.050990	0.0183
DistanceFromHome	-0.001686	-0.004985	1.000000	0.021042	NaN	0.032916	-0.0160
Education	0.208034	-0.016806	0.021042	1.000000	NaN	0.042070	-0.0271
EmployeeCount	NaN	NaN	NaN	NaN	NaN	NaN	N
EmployeeNumber	-0.010145	-0.050990	0.032916	0.042070	NaN	1.000000	0.0176
EnvironmentSatisfaction	0.010146	0.018355	-0.016075	-0.027128	NaN	0.017621	1.0000
HourlyRate	0.024287	0.023381	0.031131	0.016775	NaN	0.035179	-0.0498
JobInvolvement	0.029820	0.046135	0.008783	0.042438	NaN	-0.006888	-0.0082
JobLevel	0.509604	0.002966	0.005303	0.101589	NaN	-0.018519	0.0012
JobSatisfaction	-0.004892	0.030571	-0.003669	-0.011296	NaN	-0.046247	-0.0067
MonthlyIncome	0.497855	0.007707	-0.017014	0.094961	NaN	-0.014829	-0.0062
MonthlyRate	0.028051	-0.032182	0.027473	-0.026084	NaN	0.012648	0.0376
NumCompaniesWorked	0.299635	0.038153	-0.029251	0.126317	NaN	-0.001251	0.0125
PercentSalaryHike	0.003634	0.022704	0.040235	-0.011111	NaN	-0.012944	-0.0317
PerformanceRating	0.001904	0.000473	0.027110	-0.024539	NaN	-0.020359	-0.0295
RelationshipSatisfaction	0.053535	0.007846	0.006557	-0.009118	NaN	-0.069861	0.0076
StandardHours	NaN	NaN	NaN	NaN	NaN	NaN	N
StockOptionLevel	0.037510	0.042143	0.044872	0.018422	NaN	0.062227	0.0034
TotalWorkingYears	0.680381	0.014515	0.004628	0.148280	NaN	-0.014365	-0.0026
TrainingTimesLastYear	-0.019621	0.002453	-0.036942	-0.025100	NaN	0.023603	-0.0193
WorkLifeBalance	-0.021490	-0.037848	-0.026556	0.009819	NaN	0.010309	0.0276
YearsAtCompany	0.311309	-0.034055	0.009508	0.069114	NaN	-0.011240	0.0014
YearsInCurrentRole	0.212901	0.009932	0.018845	0.060236	NaN	-0.008416	0.0180
YearsSinceLastPromotion	0.216513	-0.033229	0.010029	0.054254	NaN	-0.009019	0.0161
YearsWithCurrManager	0.202089	-0.026363	0.014406	0.069065	NaN	-0.009197	-0.0049

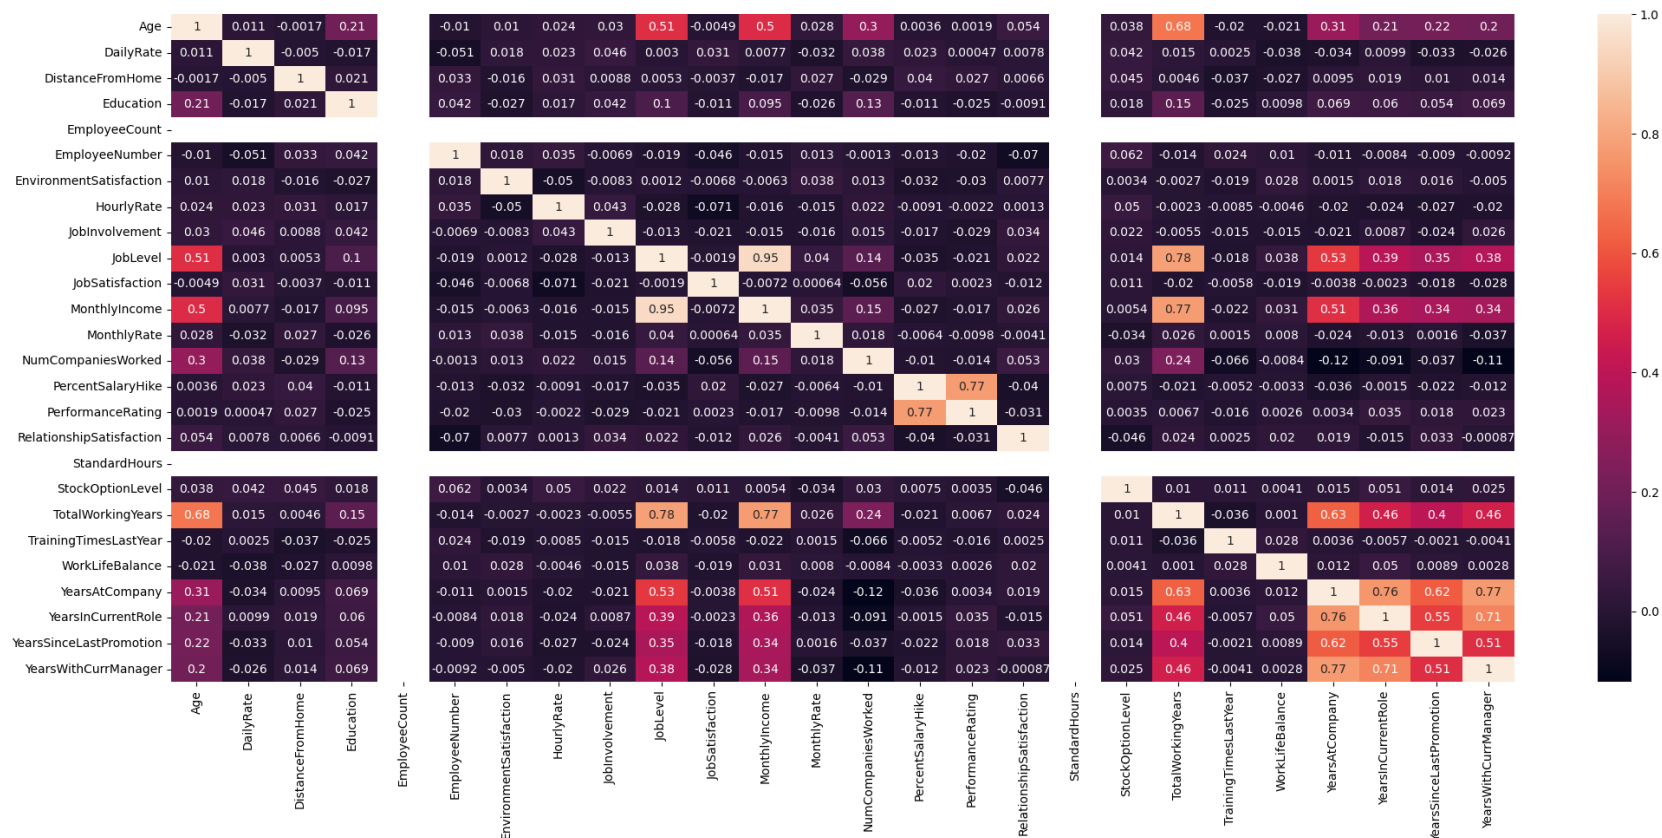
26 rows × 26 columns

```
In [13]: plt.figure(figsize=(25,10))
sns.heatmap(df.corr(),annot=True)
```

C:\Users\Praveen\AppData\Local\Temp\ipykernel_25940\1214538227.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(),annot=True)
```

Out[13]: <Axes: >



```
In [14]: len(df.columns)
```

```
Out[14]: 35
```

```
In [15]: df.head()
```

```
Out[15]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeN
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns





```
In [16]: x=df.drop("Attrition",axis=1)
x.head()
```

Out[16]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	En
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	

5 rows × 34 columns



```
In [17]: y1=df.Attrition
y=y1.to_frame()
y.head()
```

Out[17]:

	Attrition
0	Yes
1	No
2	Yes
3	No
4	No

```
In [18]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
x.BusinessTravel=le.fit_transform(x.BusinessTravel)
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)
```

```
['Non-Travel' 'Travel_Frequently' 'Travel_Rarely']
{'Non-Travel': 0, 'Travel_Frequently': 1, 'Travel_Rarely': 2}
```

```
In [19]: x['Department']=le.fit_transform(x['Department'])
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)
```

```
['Human Resources' 'Research & Development' 'Sales']
{'Human Resources': 0, 'Research & Development': 1, 'Sales': 2}
```

```
In [20]: x.EducationField=le.fit_transform(x.EducationField)
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)
```

```
['Human Resources' 'Life Sciences' 'Marketing' 'Medical' 'Other'
 'Technical Degree']
{'Human Resources': 0, 'Life Sciences': 1, 'Marketing': 2, 'Medical': 3, 'Other': 4, 'Technical Degree': 5}
```

```
In [21]: x.Gender=le.fit_transform(x.Gender)
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)
```

```
['Female' 'Male']
{'Female': 0, 'Male': 1}
```

```
In [22]: x.JobRole=le.fit_transform(x.JobRole)
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)

['Healthcare Representative' 'Human Resources' 'Laboratory Technician'
 'Manager' 'Manufacturing Director' 'Research Director'
 'Research Scientist' 'Sales Executive' 'Sales Representative']
{'Healthcare Representative': 0, 'Human Resources': 1, 'Laboratory Technician': 2, 'Manager': 3, 'Manufacturing Director': 4, 'Research Director': 5, 'Research Scientist': 6, 'Sales Executive': 7, 'Sales Representative': 8}
```

```
In [23]: x.MaritalStatus=le.fit_transform(x.MaritalStatus)
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)

['Divorced' 'Married' 'Single']
{'Divorced': 0, 'Married': 1, 'Single': 2}
```

```
In [24]: x.Over18 =le.fit_transform(x.Over18 )
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)

['Y']
{'Y': 0}
```

```
In [25]: x.OverTime =le.fit_transform(x.OverTime )
print(le.classes_)
mapping=dict(zip(le.classes_,range(len(le.classes_))))
print(mapping)

['No' 'Yes']
{'No': 0, 'Yes': 1}
```

```
In [26]: y.Attrition =le.fit_transform(y.Attrition )  
print(le.classes_)  
mapping=dict(zip(le.classes_,range(len(le.classes_))))  
print(mapping)
```

```
['No' 'Yes']  
{'No': 0, 'Yes': 1}
```

In [27]: `x.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   BusinessTravel                       1470 non-null   int32
2   DailyRate                           1470 non-null   int64
3   Department                           1470 non-null   int32
4   DistanceFromHome                    1470 non-null   int64
5   Education                           1470 non-null   int64
6   EducationField                       1470 non-null   int32
7   EmployeeCount                       1470 non-null   int64
8   EmployeeNumber                      1470 non-null   int64
9   EnvironmentSatisfaction              1470 non-null   int64
10  Gender                               1470 non-null   int32
11  HourlyRate                           1470 non-null   int64
12  JobInvolvement                       1470 non-null   int64
13  JobLevel                             1470 non-null   int64
14  JobRole                             1470 non-null   int32
15  JobSatisfaction                      1470 non-null   int64
16  MaritalStatus                       1470 non-null   int32
17  MonthlyIncome                       1470 non-null   int64
18  MonthlyRate                         1470 non-null   int64
19  NumCompaniesWorked                  1470 non-null   int64
20  Over18                              1470 non-null   int32
21  OverTime                             1470 non-null   int32
22  PercentSalaryHike                   1470 non-null   int64
23  PerformanceRating                   1470 non-null   int64
24  RelationshipSatisfaction              1470 non-null   int64
25  StandardHours                       1470 non-null   int64
26  StockOptionLevel                    1470 non-null   int64
27  TotalWorkingYears                   1470 non-null   int64
28  TrainingTimesLastYear                1470 non-null   int64
29  WorkLifeBalance                     1470 non-null   int64
30  YearsAtCompany                      1470 non-null   int64
31  YearsInCurrentRole                  1470 non-null   int64
32  YearsSinceLastPromotion              1470 non-null   int64
33  YearsWithCurrManager                 1470 non-null   int64
dtypes: int32(8), int64(26)
memory usage: 344.7 KB


```

In [28]: `x.head()`

Out[28]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Env
0	41	2	1102	2	1	2	1	1	1	
1	49	1	279	1	8	1	1	1	2	
2	37	2	1373	1	2	2	4	1	4	
3	33	1	1392	1	3	4	1	1	5	
4	27	2	591	1	2	1	3	1	7	

5 rows × 34 columns



In [29]: `y.head()`

Out[29]:

	Attrition
0	1
1	0
2	1
3	0
4	0

In [30]: `from sklearn.preprocessing import MinMaxScaler`
`ms=MinMaxScaler()`
`x_scaled=pd.DataFrame(ms.fit_transform(x),columns=x.columns)`

In [31]: `x_scaled.head()`

Out[31]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber
0	0.547619	1.0	0.715820	1.0	0.000000	0.25	0.2	0.0	0.000000
1	0.738095	0.5	0.126700	0.5	0.250000	0.00	0.2	0.0	0.000484
2	0.452381	1.0	0.909807	0.5	0.035714	0.25	0.8	0.0	0.001451
3	0.357143	0.5	0.923407	0.5	0.071429	0.75	0.2	0.0	0.001935
4	0.214286	1.0	0.350036	0.5	0.035714	0.00	0.6	0.0	0.002903

5 rows × 34 columns



In [32]: `x_scaled.shape`

Out[32]: (1470, 34)

In [33]: `y.shape`

Out[33]: (1470, 1)

In [34]: `from sklearn.model_selection import train_test_split`
`x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.2,random_state=1)`

In [35]: `print(x_train.shape)`
`print(x_test.shape)`
`print(y_train.shape)`
`print(y_test.shape)`

(1176, 34)
 (294, 34)
 (1176, 1)
 (294, 1)

Logistic Regression

```
In [36]: from sklearn.linear_model import LogisticRegression
         lr=LogisticRegression()
         lr.fit(x_train,y_train)
```

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
Out[36]: LogisticRegression
LogisticRegression()
```

```
In [37]: y_pred=lr.predict(x_test)
```

```
In [38]: y_pred
```

[illegible]

```
In [39]: y_test
```

```
Out[39]:
```

	Attrition
1291	1
1153	1
720	1
763	0
976	0
...	...
302	0
443	1
701	0
309	0
845	0

294 rows × 1 columns

```
In [40]: print(pd.DataFrame({"Actual":y_test.Attrition,"Predicted":y_pred}))
```

	Actual	Predicted
1291	1	0
1153	1	1
720	1	1
763	0	0
976	0	0
...
302	0	0
443	1	0
701	0	0
309	0	0
845	0	0

[294 rows x 2 columns]

```
In [41]: from sklearn.metrics import accuracy_score,classification_report,roc_curve
```

```
In [42]: accuracy_score(y_test,y_pred)
```

```
Out[42]: 0.8333333333333334
```


```
In [43]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.98	0.90	236
1	0.76	0.22	0.35	58
accuracy			0.83	294
macro avg	0.80	0.60	0.63	294
weighted avg	0.82	0.83	0.79	294

```
In [44]: y_probability=lr.predict_proba(x_test)[:,-1]
```

```
In [45]: fpr,tpr,thresholds=roc_curve(y_test,y_probability)
```

```
In [46]: plt.plot(fpr, tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()
```



Hyperparameter Tuning (Logistic Regression)

```
In [47]: from sklearn.model_selection import GridSearchCV

# Define a range of hyperparameters to search
param_grid = {
    'penalty': ['l1', 'l2'],
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Inverse of regularization strength
    'solver': ['liblinear', 'saga']
}

# Create a Logistic regression classifier
lr = LogisticRegression()

# Create a GridSearchCV object with 5-fold cross-validation
grid_search = GridSearchCV(lr, param_grid, cv=5, scoring='accuracy')

# Fit the grid search to your training data
grid_search.fit(x_train, y_train)
```

column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\linear_model_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

warnings.warn(

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\linear_model_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

warnings.warn(

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
In [49]: # Print the best hyperparameters found by grid search
print("Best Hyperparameters:")
print(grid_search.best_params_)
```

```
Best Hyperparameters:
{'C': 10, 'penalty': 'l1', 'solver': 'saga'}
```

```
In [50]: # Get the best model
best_lr = grid_search.best_estimator_
best_lr
```

```
Out[50]: LogisticRegression
LogisticRegression(C=10, penalty='l1', solver='saga')
```

```
In [51]: # Evaluate the best model on the test set
y_pred_best = best_lr.predict(x_test)
print(pd.DataFrame({"Actual":y_test.Attrition,"Predicted":y_pred_best}))
```

	Actual	Predicted
1291	1	0
1153	1	1
720	1	1
763	0	0
976	0	0
...
302	0	0
443	1	0
701	0	0
309	0	0
845	0	0

[294 rows x 2 columns]

```
In [52]: # Calculate accuracy and other metrics
accuracy_best = accuracy_score(y_test, y_pred_best)
classification_report_best = classification_report(y_test, y_pred_best)

print("Accuracy with Best Hyperparameters:", accuracy_best)
print("Classification Report with Best Hyperparameters:\n", classification_report_best)
```

Accuracy with Best Hyperparameters: 0.8469387755102041

Classification Report with Best Hyperparameters:

	precision	recall	f1-score	support
0	0.85	0.98	0.91	236
1	0.78	0.31	0.44	58
accuracy			0.85	294
macro avg	0.82	0.64	0.68	294
weighted avg	0.84	0.85	0.82	294

Decision Tree

```
In [53]: from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
```

```
In [54]: dtc.fit(x_train,y_train)
```

```
Out[54]: ▾ DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [55]: y_pred=dtc.predict(x_test)
```

```
In [56]: print(pd.DataFrame({"Actual":y_test.Attrition,"Predicted":y_pred}))
```

	Actual	Predicted
1291	1	1
1153	1	1
720	1	0
763	0	0
976	0	1
...
302	0	1
443	1	0
701	0	0
309	0	0
845	0	0

[294 rows x 2 columns]


```
In [57]: accuracy_score(y_test,y_pred)
```

Out[57]: 0.7789115646258503

```
In [58]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.89	0.87	236
1	0.42	0.31	0.36	58
accuracy			0.78	294
macro avg	0.63	0.60	0.61	294
weighted avg	0.76	0.78	0.77	294


```
In [59]: y_prob=dtc.predict_proba(x_test)[:,-1]
fpr,tpr,thresholds=roc_curve(y_test,y_prob)
plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()
```



Hyperparameter tuning(Pre Pruning) for Decision Tree

```
In [60]: para={
    'criterion':['entropy','gini'],
    'splitter':['best','random'],
    'max_features':['auto','sqrt','log2'],
    'max_depth':list(range(0,10))
}
```

```
In [61]: grid_dtc=GridSearchCV(dtc,para,cv=10,scoring='accuracy')
```

```
In [62]: grid_dtc.fit(x_train,y_train)
```

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\tree_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.

warnings.warn(

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\tree_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.

warnings.warn(

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\tree_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.

warnings.warn(

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\tree_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.

warnings.warn(

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\tree_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.

.....

```
In [63]: grid_dtc.best_params_
```

```
Out[63]: {'criterion': 'entropy',
          'max_depth': 7,
          'max_features': 'auto',
          'splitter': 'random'}
```

```
In [64]: best_dtc=DecisionTreeClassifier(
          criterion=grid_dtc.best_params_['criterion'],
          max_depth=grid_dtc.best_params_['max_depth'],
          max_features=grid_dtc.best_params_["max_features"],
          splitter=grid_dtc.best_params_["splitter"]
        )
        best_dtc.fit(x_train,y_train)
```

C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\tree_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.

warnings.warn(

```
Out[64]: DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=7, max_features='auto',
                      splitter='random')
```

```
In [65]: y_pred=best_dtc.predict(x_test)
          print(pd.DataFrame({"Actual":y_test.Attrition,"Predicted":y_pred}))
```

	Actual	Predicted
1291	1	0
1153	1	0
720	1	0
763	0	0
976	0	0
...
302	0	0
443	1	0
701	0	0
309	0	0
845	0	0

[294 rows x 2 columns]


```
In [66]: print(accuracy_score(y_test,y_pred))
print(classification_report(y_test,y_pred))
```

```
0.7993197278911565
      precision    recall  f1-score   support

     0       0.82      0.95      0.88       236
     1       0.48      0.17      0.25        58

 accuracy          0.80       294
 macro avg       0.65      0.56      0.57       294
 weighted avg    0.76      0.80      0.76       294
```

```
In [88]: from sklearn import tree
plt.figure(figsize=(80,60))
tree.plot_tree(best_dtc,filled=True)
```



Random Forest

```
In [67]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
```

```
In [68]: rfc.fit(x_train,y_train)
```

C:\Users\Praveen\AppData\Local\Temp\ipykernel_25940\4070307935.py:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
rfc.fit(x_train,y_train)
```

```
Out[68]: RandomForestClassifier
RandomForestClassifier()
```

```
In [69]: y_pred=rfc.predict(x_test)
```

```
In [70]: accuracy_score(y_test,y_pred)
```

```
Out[70]: 0.8231292517006803
```

```
In [71]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.83	0.99	0.90	236
1	0.75	0.16	0.26	58
accuracy			0.82	294
macro avg	0.79	0.57	0.58	294
weighted avg	0.81	0.82	0.77	294

Pre pruning Random Forest

```
In [92]: para={
    'criterion':['gini','entropy'],
    'max_features':['best','sqrt','log2',None],
    'max_depth':[10, 20, 30, None],
}
```

```
In [93]: rfc_cv=GridSearchCV(rfc,para,cv=5,scoring="accuracy")
```

```
In [94]: rfc_cv.fit(x_train,y_train)
```

```
les,)), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,)), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,)), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,)), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
C:\Users\Praveen\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:686: DataConversionWarning: A column-vector v was passed when a 1d array was expected. Please change the shape of v to (n_samp
```

```
In [95]: rfc_cv.best_params_
```

```
Out[95]: {'criterion': 'entropy', 'max_depth': 10, 'max_features': None}
```

```
In [96]: best_rfc=RandomForestClassifier(
        criterion=rfc_cv.best_params_['criterion'],
        max_depth=rfc_cv.best_params_['max_depth'],
        max_features=rfc_cv.best_params_["max_features"],
    )
```

```
In [97]: best_rfc.fit(x_train,y_train)
```

C:\Users\Praveen\AppData\Local\Temp\ipykernel_25940\2820291153.py:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
best_rfc.fit(x_train,y_train)
```

```
Out[97]: RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=10, max_features=None)
```

```
In [98]: y_pred=best_rfc.predict(x_test)
```

```
In [101]: print("Accuracy:",accuracy_score(y_test,y_pred))
print("Classification Report:\n",classification_report(y_test,y_pred))
```

Accuracy: 0.8231292517006803

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.98	0.90	236
1	0.69	0.19	0.30	58
accuracy			0.82	294
macro avg	0.76	0.58	0.60	294
weighted avg	0.80	0.82	0.78	294

```
In [103]: y_prob=best_rfc.predict_proba(x_test)[:,-1]
```

```
In [106]: fpr, tpr, thresholds = roc_curve(y_test, y_prob)
plt.plot(fpr, tpr)
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC CURVE")
plt.show()
```

